


A Bayesian Optimization Approach to Machine Translation Reranking

Julius Cheng,
Maike Züfle,
Vilém Zouhar,
Andreas Vlachos



Problem: MT produces 100 translations; how to find the best one?

Idea 0: Use MT logprobs to select the most probable one. \rightsquigarrow low quality

Idea 1: Use quality estimation metric. \rightsquigarrow expensive

Idea 2: Don't evaluate all candidates.

Estimate unknown scores with existing scores or cheap proxy metrics.

Idea 3: Search efficiently with Gaussian Process and BayesOpt.

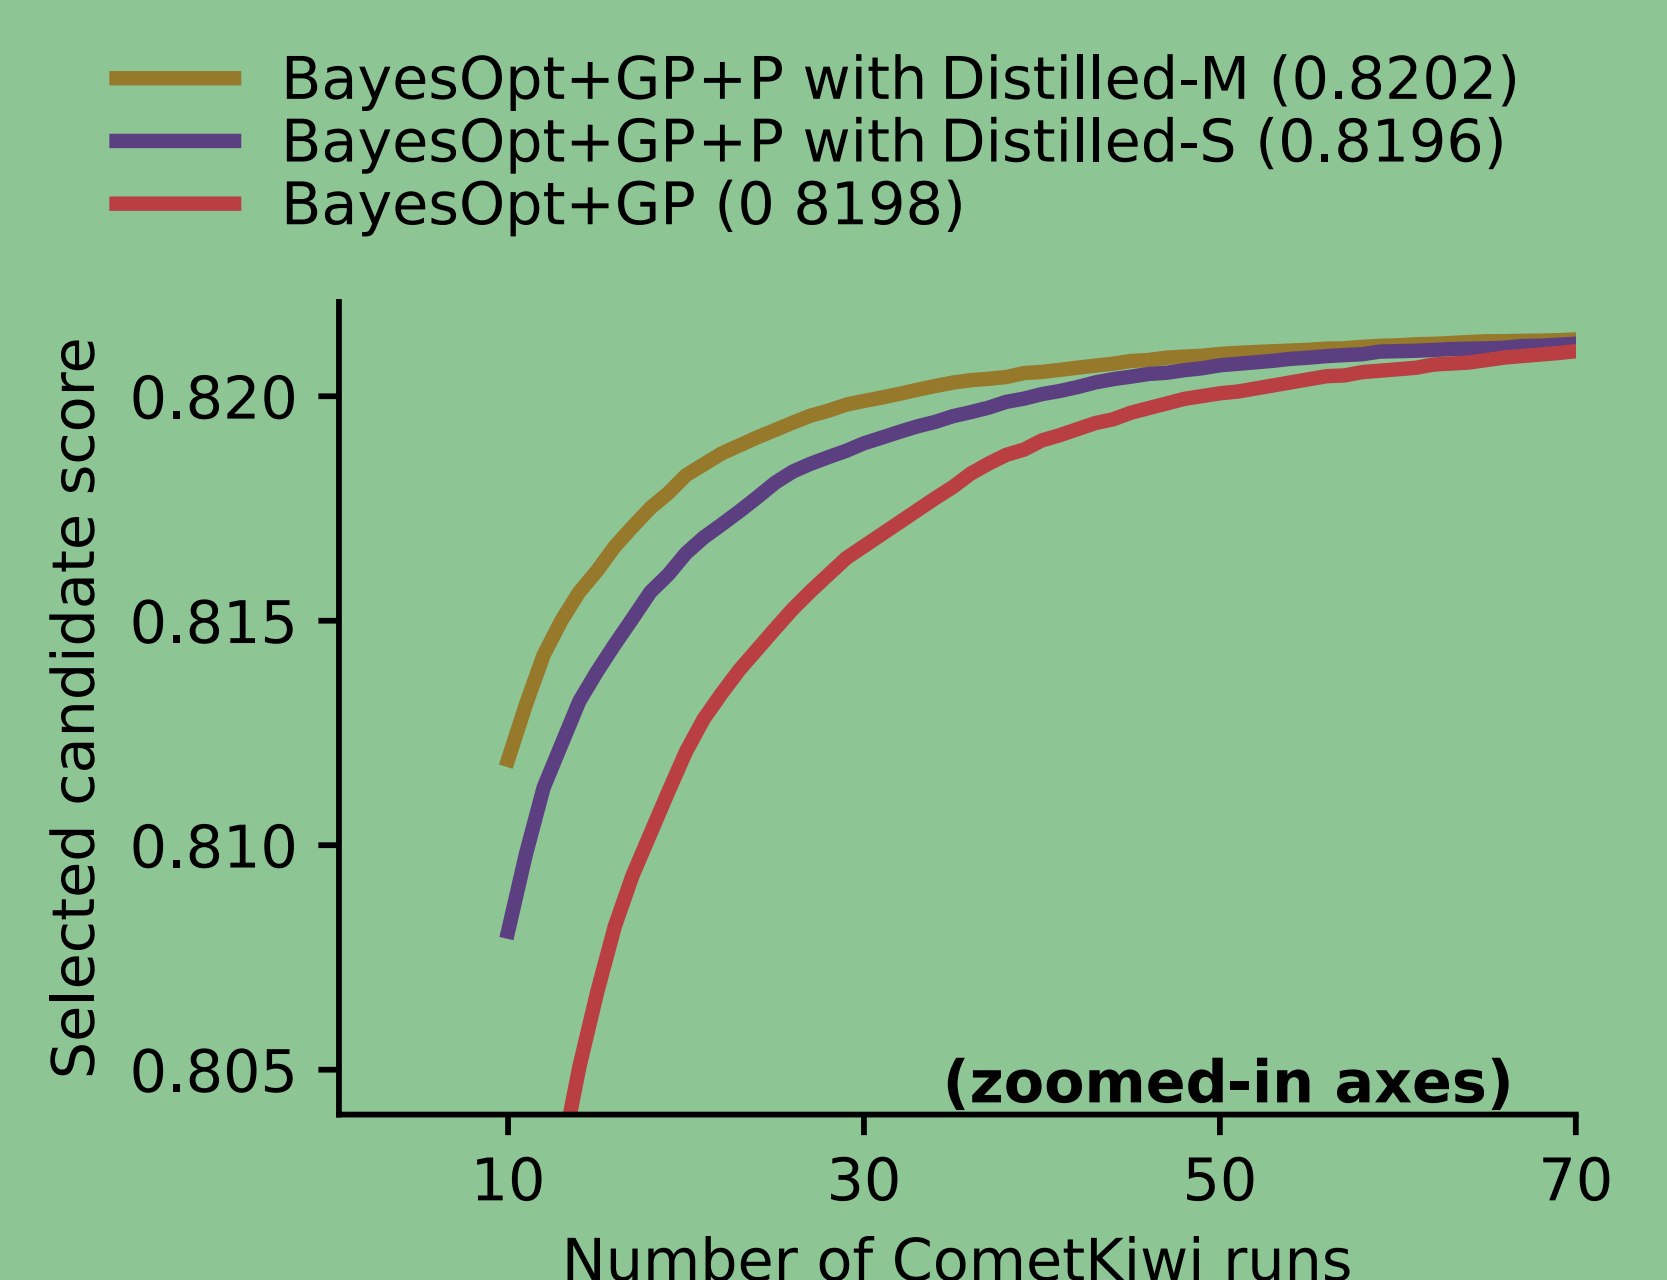

```
1: while  $|C_{\text{seen}}| < \text{Budget}$  :
2:    $y_{\text{next}} = \arg \max_{y \in C - C_{\text{seen}}} \text{ExpectedImprovement}(y, C_{\text{seen}})$ 
3:    $C_{\text{seen}} = C_{\text{seen}} \cup \{y_{\text{next}}\}$  and compute  $\text{score}(y_{\text{next}})$ 
4: return  $\arg \max_{y \in C_{\text{seen}}} \text{score}(y)$ 
```

Iteratively select the most promising candidate to score. Return the best one seen so far.

Compute expected improvement based on scores revealed so far for other candidates and their distance to y

```
1:  $\mathcal{K}(y, y') = \exp\left(-\frac{\|\text{embd}(y) - \text{embd}(y')\|^2}{2w^2}\right)$ 
2: def ExpectedImprovement( $y, C_{\text{seen}}$ ) :
3:    $y_{\text{best}} = \arg \max_{y \in C_{\text{seen}}} \text{score}(y)$ 
4:    $\sigma(y), \mu(y) = \text{GaussianProcess}(\mathcal{K}, y, C_{\text{seen}})$ 
5:    $z = (\text{score}(y_{\text{best}}) - \mu(y)) / \sigma(y)$ 
6: return  $\sigma(y)(z \cdot \text{cdf}(z) + \text{pdf}(z))$ 
```

Find near-best candidates with much fewer scorings



Baselines:

UniqRandom: score B random candidates & select max.

Logprob: score top-B candidates based on logprob.

HillClimbing: iteratively score closest to current best.

BayesOpt+GP: our algorithm

Upgrade:

Bootstrap selection with cheaper but less accurate scorers (small COMET quality estimation).